

CMU 10-715: Homework 6 Report
Implementation of Convolutional Neural Networks
Abishek Sridhar (Andrew Id: abisheks).

1 Results

- a I trained the classifiers for **3000** steps on Google Colab's GPU.
- b Refer table 1 for the final results of the training.

Model	Pooling	Training Time	Train Acc	Test Acc	Steps
M_{max}	max	518.707s	79.900	64.080	3000
M_{avg}	avg	515.660s	75.672	62.380	3000
M_{no}	no	528.855s	99.936	60.360	3000

Table 1: Table showing the final training and testing accuracies, total training time (including evaluation), and training steps for the three models with different pooling configurations

We can observe from the table that models M_{avg} and M_{max} have similar training times which is markedly less than that of M_{no} . This is because of removal of the subsampling pooling layers. In the CNN architecture under consideration, the **Conv2d** layers have a stride of 1 and do not result in much shrinkage to the feature maps's size. It is the pooling layers (**max** or **avg** configurations) with a stride of 2 that result in shrinkage of the feature maps' dimensions across the layers as the number of feature maps grow. Consequently, the first fully connected layer that comes up after the **nn.Flatten()** have smaller input dimensions in the case of M_{avg} and M_{max} (numerically, = 400) compared to M_{no} (input dim for first FC = 9216). Hence, there are far more parameters and bigger computation graph (for back-propagation) in the no pooling case (the pooling layers have no trainable parameters). This overparameterization is also evident in the way M_{no} overfits to the training data, as can be seen in table 1. Though there are additional computations in the pooling layers in M_{avg} and M_{max} models, it is seemingly outweighed by the computational time of M_{no} from its additional parameters, computational graph nodes, and back-propagation updates.

- c Refer figure 1.
- d Refer figure 2.
- e Refer figure 3.

Let's first compare the output of conv1 to the input image. We see that the six output feature maps of conv1 layer all capture/highlight varied

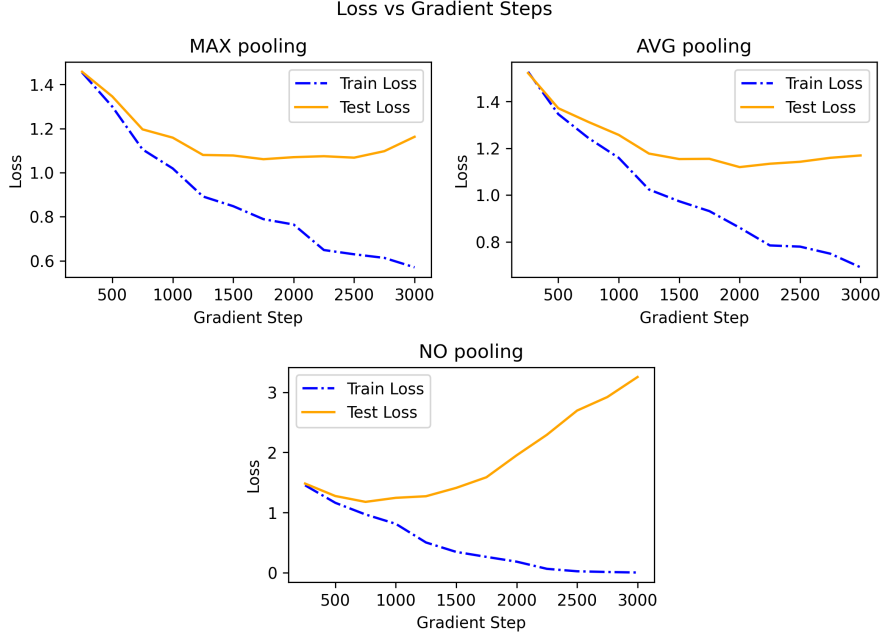


Figure 1: Loss vs Gradient Step plots for models M_{max} , M_{avg} , and M_{no} till 3000 training steps

features in the same input image, but they are truncated on the sides (we can notice missing pixels on the edges due to the convolution with no padding). In each of the feature maps, we can observe some characteristic edge or feature of the frog being highlighted and others being blurred out.

Now, let's compare the output of pool1 to the output of conv1. Clearly, the six images are now more pixelated due to the subsampling pooling layer that reduced the width and height dimensions by half. We can still observe that the features/edges that were highlighted in the six features maps after conv1 respectively seem to be preserved in the six feature maps after pool1 layer, but they are smoothened out and not sharp now. This might be due to the usage of average pooling - the boundaries of the features identified by conv1 are no longer uniform/sharp but contain mixed pixel values due to being affected by neighboring pixel values.

To sum it up, the convolution layers seem to be identifying features in the input image that it considers predictive of the output labels and the pooling layers seem to perform the job of reducing their dimensions for reducing overparameterization and computation cost, preserving as much

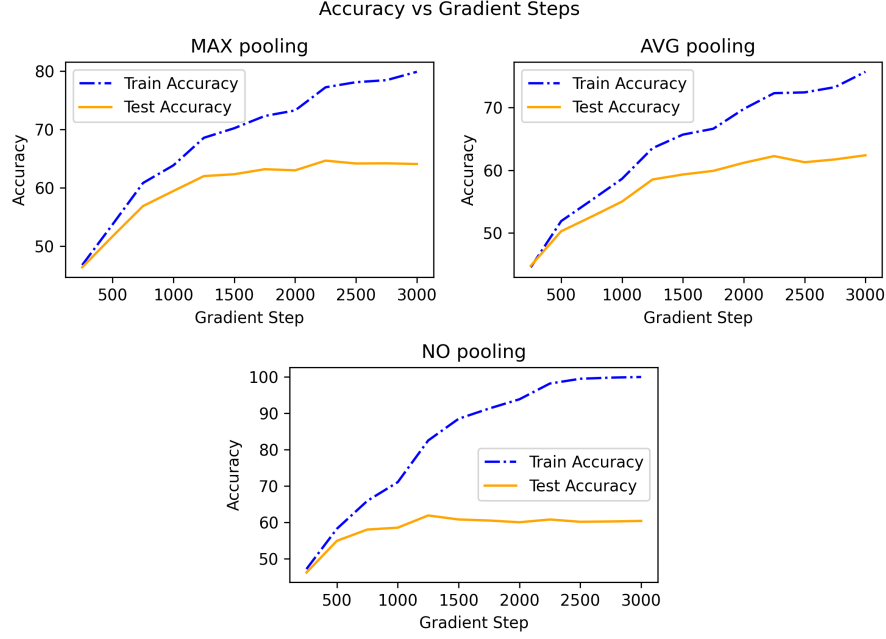


Figure 2: Accuracy vs Gradient Step plots for models M_{max} , M_{avg} , and M_{no} till 3000 training steps

information extracted by the convolutional layers as possible. Max pooling might be doing better than average pooling (as seen in table 1) because it will preserve the values of the important pixels without distorting them, thereby preserving information in a better manner than average pooling.

- f The architecture of the convolutional neural network that I designed (M_{custom}) is as follows:

Input image \rightarrow 2D convolution with output channels = 8, kernel size 5, stride 1 and padding = 'same' \rightarrow ReLU \rightarrow Max Pooling with kernel size 2, stride 2 and no padding \rightarrow 2D convolution with output channels = 16, kernel size 5, stride 1 and padding = 'same' \rightarrow ReLU \rightarrow Max Pooling with kernel size 2, stride 2 and no padding \rightarrow 2D convolution with output channels = 64, kernel size 3, stride 1 and padding = 'same' \rightarrow ReLU \rightarrow Max Pooling with kernel size 2, stride 2 and no padding \rightarrow 2D convolution with output channels = 128, kernel size 4, stride 1 and no padding \rightarrow ReLU \rightarrow Fully connected linear layer with input dim = 128, output dim = 80 \rightarrow ReLU \rightarrow Fully connected linear layer with input dim = 80, output dim = 10 \rightarrow Output

The intuitions behind the modifications to LeNet architecture were drawn from the various observations during this assignment: used padding = 'same' for the convolutional layers for it to see the feature maps completely and not lose information at the sides of the feature maps (since images did extend till the sides); tried to create a slightly deeper network hence added an additional convolutional layer; to fight over-parameterization replaced the first FC layer by a convolutional layer and used subsampling layer after every convolutional layer; used max pooling since it performed better than average pooling by extracting extreme features well; finally, gradually increased number of output channels of convolutional layers as the feature map dimensions halved after the pooling layers.

I trained M_{custom} for the same number of training steps as LeNet for a fair comparison and obtained the final results as summarized in 2. The custom model does contain an additional layer and hence more parameters, that leads to the increased training set accuracy. But, the testing set accuracy (\Rightarrow generalization performance) also increases by a non-trivial amount and validates the intuitions behind the architecture. The overall training time taken by M_{custom} is higher than M_{max} as expected due to deeper network, and just lesser than M_{no} , but the improved performance compensates for this marginal increase in training time.

M_{custom} actually achieves highest its best performance at ≈ 2250 steps, with a test accuracy of 70.940 and train accuracy of 91.222, which can be obtained by early stopping the training.

Model	Pooling	Training Time	Train Acc	Test Acc	Steps
M_{max}	max	518.707s	79.900	64.080	3000
M_{custom}	max	527.135s	94.538	69.340	3000


Table 2: Table showing the final training and testing accuracies, total training time (including evaluation), and training steps for the best LeNet model (M_{max}) and custom architected CNN model

Input Image




(a) The original image - one 32×32 image

Output Feature Maps of conv1



(b) Output of the first layer (after the first convolutional layer but before the ReLU)
- six 28×28 images

Output Feature Maps of pool1



(c) Output of the second layer (after the first pooling layer and before the second
convolutional layer) - six 14×14 images

Figure 3: Visualization of input image and initial layer outputs for the input image

2 Code

Colab_Cell_1

```
1 from google.colab import drive
2 drive.mount('/content/drive')
3
4 %cd /content/drive/MyDrive/10-715/HW6
5 !mkdir -p ./results ./plots ./models
```

Colab_Cell_2

```
1 %matplotlib inline
2
3 import time
4 import numpy as np
5 import matplotlib.pyplot as plt
6 import torch as t
7 from torch import nn
8 from torch import optim
9 from torch.nn import functional as F
10 import torchvision
11 import torchvision.transforms as transforms
12 import pickle
13 import itertools
14
15
16 def divide_no_nan(a, b):
17     """
18     a/b where the resulted NaN or Inf are replaced by 0.
19     """
20     result = a / b
21     result[result != result] = .0
22     result[result == np.inf] = .0
23     return result
24
25
26 class WeightNormConstrainer(object):
27     def __init__(self, norm):
28         self.norm = norm
29
30     def __call__(self, module):
31         if hasattr(module, 'weight'):
32             w = module.weight.data
```

```

33         wn = t.norm(w, p=2, dim=1).detach()
34         ind = t.gt(wn, self.norm)
35         div = (divide_no_nan(wn, self.norm) * ind) + (1 *
36             ↪ t.logical_not(ind))
37         div = div.unsqueeze_(1)
38         w.div_(div)
39
40 class _LeNet(nn.Module):
41     def __init__(self, layers):
42         super(_LeNet, self).__init__()
43         self.input_layer = nn.Sequential(*layers['input_layer'])
44         self.hidden_layers = nn.Sequential(*layers['hidden_layers'])
45         self.output_layer = nn.Sequential(*layers['output_layer'])
46         self.activation = {}
47
48     def layer_features(self, x):
49         input_layer = self.input_layer(x)
50         hidden = self.hidden_layers(input_layer)
51         logits = self.output_layer(hidden)
52         layer_features = {'input_layer': input_layer.data.cpu().numpy(),
53             ↪ 'hidden': hidden.data.cpu().numpy(),
54             ↪ 'logits': logits.data.cpu().numpy()}
55         return layer_features
56
57     def get_activation(self, name):
58         def hook(model, input, output):
59             self.activation[name] = output.detach().cpu().numpy()
60         return hook
61
62     def get_intermediate(self, x):
63         #TODO: complete this function to return the outputs of first conv and
64         ↪ first pooling layer - Done
65         # Returns the output as a numpy array for ease of plotting it later
66         self.input_layer[1].register_forward_hook(self.get_activation('conv1')
67             ↪ ))
68         pool1_out = self.input_layer(x)
69         return self.activation['conv1'], pool1_out.detach().cpu().numpy()
70
71     def forward(self, x):
72         input_layer = self.input_layer(x)
73         hidden = self.hidden_layers(input_layer)
74         logits = self.output_layer(hidden)
75         return logits

```



```

73 # We recommend you make this modules as the first component of your input
    ↪ layers
74 class Reshape(t.nn.Module):
75     def forward(self, x):
76         return x.view(-1,3,32,32)
77
78 class LeNet(object):
79     def __init__(self, params, use_custom=False):
80         super().__init__()
81         self.params = params
82         self.device = t.device('cuda' if t.cuda.is_available() else 'cpu')
83         self.activations = {'logistic': nn.Sigmoid(), 'relu': nn.ReLU()}
84         self.pooling = {'avg':nn.AvgPool2d(kernel_size=2,
    ↪ stride=2), 'max':nn.MaxPool2d(kernel_size=2, stride=2)}
85         # Instantiate model
86         t.manual_seed(self.params['random_seed'])
87         np.random.seed(self.params['random_seed'])
88         layers = self._initialize_network() if not use_custom else
    ↪ self._initialize_network_custom()
89         self.model = _LeNet(layers).to(self.device)
90
91     def _initialize_network(self):
92         # TODO: complete this function - Done
93         # You may have to use Reshape() and nn.Flatten()
94         # The architecture may be slightly different with no pooling
95         if self.params['pooling'] != 'no':
96             input_layers = [
97                 Reshape(),
98                 nn.Conv2d(3, 6, kernel_size=5, stride=1, padding=0),
99                 self.activations[self.params['activation']],
100                 self.pooling[self.params['pooling']]
101             ]
102             hidden_layers = [
103                 nn.Conv2d(6, 16, kernel_size=5, stride=1, padding=0),
104                 self.activations[self.params['activation']],
105                 self.pooling[self.params['pooling']],
106                 # nn.Conv2d(16, 120, kernel_size=5, stride=1, padding=0),
107                 # self.activations[self.params['activation']],
108                 nn.Flatten(),
109                 nn.Linear(400, 120),
110                 self.activations[self.params['activation']],
111                 nn.Linear(120, 84),
112                 self.activations[self.params['activation']]
113             ]

```

```

114         output_layer = [
115             nn.Linear(84, 10)
116         ]
117     else:
118         input_layers = [
119             Reshape(),
120             nn.Conv2d(3, 6, kernel_size=5, stride=1, padding=0),
121             self.activations[self.params['activation']]
122         ]
123         hidden_layers = [
124             nn.Conv2d(6, 16, kernel_size=5, stride=1, padding=0),
125             self.activations[self.params['activation']],
126             # nn.Conv2d(16, 120, kernel_size=5, stride=1, padding=0),
127             # self.activations[self.params['activation']],
128             nn.Flatten(),
129             nn.Linear(9216, 120),
130             self.activations[self.params['activation']],
131             nn.Linear(120, 84),
132             self.activations[self.params['activation']]
133         ]
134         output_layer = [
135             nn.Linear(84, 10)
136         ]
137     layers = {'input_layer': input_layers,
138             'hidden_layers': hidden_layers,
139             'output_layer': output_layer}
140     return layers
141
142 def _initialize_network_custom(self):
143     # Custom CNN network
144     assert self.params['pooling'] != 'no'
145
146     input_layers = [
147         Reshape(),
148         nn.Conv2d(3, 8, kernel_size=5, stride=1, padding='same'),
149         self.activations[self.params['activation']],
150         self.pooling[self.params['pooling']]
151     ]
152     hidden_layers = [
153         nn.Conv2d(8, 16, kernel_size=5, stride=1, padding='same'),
154         self.activations[self.params['activation']],
155         self.pooling[self.params['pooling']],
156         nn.Conv2d(16, 64, kernel_size=3, stride=1, padding='same'),
157         self.activations[self.params['activation']],

```

```

158         self.pooling[self.params['pooling']],
159         nn.Conv2d(64, 128, kernel_size=4, stride=1, padding=0),
160         self.activations[self.params['activation']],
161         nn.Flatten(),
162         nn.Linear(128, 80),
163         self.activations[self.params['activation']]
164     ]
165     output_layer = [
166         nn.Linear(80, 10)
167     ]
168
169     layers = {'input_layer': input_layers,
170             'hidden_layers': hidden_layers,
171             'output_layer': output_layer}
172     return layers
173
174
175 def evaluate_accuracy(self, dataloader):
176     self.model.eval()
177     with t.no_grad():
178         n_correct = 0
179         n_samples = 0
180         for batch in iter(dataloader):
181             batch_x = t.flatten(batch[0].to(self.device), start_dim=1)
182             batch_y = batch[1].to(self.device)
183             logits = self.model(batch_x)
184             y_hat = t.argmax(logits, dim=1)
185             correct = t.sum(y_hat==batch_y)
186             n_correct += correct.data.cpu().numpy()
187             n_samples += len(batch_x)
188         accuracy = (n_correct/n_samples) * 100
189         self.model.train()
190         return accuracy
191
192 def evaluate_cross_entropy(self, dataloader):
193     self.model.eval()
194     loss_func = nn.CrossEntropyLoss()
195     with t.no_grad():
196         n_samples = 0
197         total_loss = 0
198         for batch in iter(dataloader):
199             pass
200             # TODO: Write the code to measure the cross_entropy
201             # (Hint, look at the evaluate_accuracy method)

```

```

202         # Be careful with the eval and train modes of the model
203         # This should be same as HW5 - Done
204         batch_x = t.flatten(batch[0].to(self.device), start_dim=1)
205         batch_y = batch[1].to(self.device)
206
207         logits = self.model(batch_x)
208         loss = loss_func(logits, batch_y)
209
210         total_loss += loss.item() * len(batch_x)
211         n_samples += len(batch_x)
212
213     cross_entropy = total_loss/n_samples
214
215     self.model.train()
216
217     return cross_entropy
218
219 def adjust_lr(self, optimizer, lr_decay):
220     for param_group in optimizer.param_groups:
221         param_group['lr'] = param_group['lr'] * lr_decay
222
223 def adjust_momentum(self, optimizer, step, momentum_change_steps,
224                     initial_momentum, final_momentum):
225     mcs = momentum_change_steps
226     s = min(step, mcs)
227     momentum = (initial_momentum) * ((mcs-s)/mcs) + \
228               (final_momentum) * (s/mcs)
229     for param_group in optimizer.param_groups:
230         param_group['momentum'] = momentum
231
232 def save_weights(self, path):
233     t.save(self.model.state_dict(), path)
234
235 def load_weights(self, path):
236     self.model.load_state_dict(t.load(path,
237                                     map_location=t.device(self.device)))
238     self.model.eval()
239
240 def fit(self, insample_dataloader, outsample_dataloader):
241     # Instantiate optimization tools
242     loss = nn.CrossEntropyLoss()
243     optimizer = optim.SGD([{'params':
        ↪ self.model.input_layer.parameters()}],

```

```

244         {'params':
           ↪ self.model.hidden_layers.parameters()},
245         {'params':
           ↪ self.model.output_layer.parameters(),
246          'weight_decay':
           ↪ self.params['output_l2_decay']}},
247         lr=self.params['initial_lr'],
248         momentum=self.params['initial_momentum'])
249
250     constrainer = WeightNormConstrainer(norm=self.params['weight_norm'])
251     # Initialize counters and trajectories
252     step = 0
253     epoch = 0
254     metric_trajectories = {'step': [],
255                            'epoch': [],
256                            'insample_accuracy': [],
257                            'outsample_accuracy': [],
258                            'insample_cross_entropy': [],
259                            'outsample_cross_entropy': []
260                           }
261
262     print('\n'+ '='*36+f' Fitting LeNet ({self.params["pooling"]} pooling)
           ↪ '+'*36)
263     while step <= self.params['iterations']:
264
265         # Train
266         epoch += 1
267         start_time = time.time()
268         self.model.train()
269         for batch in iter(insample_dataloader):
270             step+=1
271             if step > self.params['iterations']:
272                 continue
273
274             batch_x = t.flatten(batch[0].to(self.device), start_dim=1)
275             batch_y = batch[1].to(self.device)
276
277             optimizer.zero_grad()
278
279             # TODO: make predictions, compute the cross entropy loss and
280             ↪ perform backward propagation
281             # This should be same as HW5 - Done
282             logits = self.model(batch_x)
283             cross_entropy = loss(logits, batch_y)

```

```

283         cross_entropy.backward()
284
285         t.nn.utils.clip_grad_norm_(self.model.parameters(), 20)
286         optimizer.step()
287
288         # Evaluate metrics
289         if (step % self.params['display_step'] == 0):
290             in_cross_entropy =
291                 ↪ self.evaluate_cross_entropy(insample_dataloader)
292             out_cross_entropy =
293                 ↪ self.evaluate_cross_entropy(outsample_dataloader)
294             in_accuracy =
295                 ↪ self.evaluate_accuracy(insample_dataloader)
296             out_accuracy =
297                 ↪ self.evaluate_accuracy(outsample_dataloader)
298
299             print('Epoch:', '%d,' % epoch,
300                   'Step:', '%d,' % step,
301                   'In Loss: {:.7f}'.format(in_cross_entropy),
302                   'Out Loss: {:.7f}'.format(out_cross_entropy),
303                   'In Acc: {:.3f}'.format(in_accuracy),
304                   'Out Acc: {:.3f}'.format(out_accuracy))
305
306             metric_trajectories['insample_cross_entropy'].append(in_c
307                 ↪ ross_entropy)
308             metric_trajectories['outsample_cross_entropy'].append(out
309                 ↪ _cross_entropy)
310             metric_trajectories['insample_accuracy'].append(in_accura
311                 ↪ cy)
312             metric_trajectories['outsample_accuracy'].append(out_accu
313                 ↪ racy)
314             metric_trajectories['step'].append(step)
315             metric_trajectories['epoch'].append(epoch)
316
317         # Update optimizer learning rate
318         if step % self.params['adjust_lr_step'] == 0:
319             self.adjust_lr(optimizer=optimizer,
320                 ↪ lr_decay=self.params['lr_decay'])
321
322         # Update optimizer momentum
323         if step % self.params['adjust_momentum_step'] == 0 and \
324             step < self.params['momentum_change_steps']:
325             self.adjust_momentum(optimizer=optimizer, step=step,

```

```

317         momentum_change_steps=self.params['m_
    ↪ omentum_change_steps'],
318         initial_momentum=self.params['initia_
    ↪ l_momentum'],
319         final_momentum=self.params['final_mo_
    ↪ mentum'])

320
321         # Constraint max_norm of weights
322         if self.params['apply_weight_norm'] and \
323             (step % self.params['adjust_norm_step'] == 0):
324             self.model.apply(constrainer)
325
326         print(f"Time elapsed for epoch {epoch}: {time.time()-start_time}
    ↪ seconds")
327
328         # Store trajectories
329         print('\n'+ '='*35+ ' Finished Train '+'='*35)
330         self.trajectories = metric_trajectories
331
332
333     def main(POOLING, custom=False):
334         params = {'model': 'lenet',
335                   'display_step': 250,
336                   'batch_size': 256,
337                   'iterations': 3000, # 3000 or 500
338                   'initial_lr': 0.05,
339                   'lr_decay': 0.5,
340                   'adjust_lr_step': 1000,
341                   'initial_momentum': 0.9,
342                   'final_momentum': 0.95,
343                   'momentum_change_steps': 5_000,
344                   'adjust_momentum_step': 2_000,
345                   'apply_weight_norm': True,
346                   'weight_norm': 3.5,
347                   'adjust_norm_step': 1_000,
348                   'output_l2_decay': 0.001,
349                   'pooling': POOLING, # TODO try three options: 'max', 'avg', 'no'
350                   'activation': 'relu',
351                   'random_seed': 0}
352
353         transform = transforms.Compose([transforms.ToTensor(),
    ↪ transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
354         trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
    ↪ download=True, transform=transform)

```

```

355     insample_dataloader = t.utils.data.DataLoader(trainset,
356     ↪     batch_size=params['batch_size'], shuffle=True)
357     testset = torchvision.datasets.CIFAR10(root='./data', train=False,
358     ↪     download=True, transform=transform)
359     outsample_dataloader = t.utils.data.DataLoader(testset,
360     ↪     batch_size=params['batch_size'], shuffle=False)
361
362     clf = LeNet(params, custom)
363
364     t0 = time.time()
365     clf.fit(insample_dataloader, outsample_dataloader)
366     print(f"Total time elapsed for training: {time.time()-t0} seconds")
367
368     # TODO: save trajectories and/or plot trajectories, final accuracy, and
369     ↪ time elapsed - Done
370     # TODO: plot the intermediate outputs using get_intermediate in _LeNet
371     # To avoid unnecessary pain, saving classifier weights
372     if not custom:
373         clf.save_weights(f'./models/model_{POOLING}.pth')
374         clf.load_weights(f'./models/model_{POOLING}.pth')
375
376         with open(f'./results/results_{POOLING}.pkl', 'wb') as file:
377             pickle.dump(clf.trajectories, file)
378
379         # Pick one sample from the training set (so that its easier to
380         ↪ explain what the model is learning)
381         index = 1
382         batch_x, batch_y = next(itertools.islice(iter(insample_dataloader),
383         ↪     index, None))
384         sample_x, sample_y = t.flatten(batch_x[:1].to(clf.device),
385         ↪     start_dim=1), batch_y[0].item()
386         conv1_out, pool1_out = clf.model.get_intermediate(sample_x)
387         sample_x = sample_x.detach().cpu().numpy()
388         intermediates = {
389             'input': sample_x,
390             'label': sample_y,
391             'conv1': conv1_out,
392             'pool1': pool1_out
393         }
394
395         with open(f'./results/visualize_{POOLING}.pkl', 'wb') as file:
396             pickle.dump(intermediates, file)
397
398     del clf

```



```

392
393 if __name__ == '__main__':
394     for POOLING in ['max']:
395         main(POOLING, True) # The second argument is True to train custom
                               ↪ architecture; False for training LeNet architecture

```

Colab_Cell_3

```

1 import gc
2 import torch
3
4 # Code to clear any memory trace left by previous models to get proper
  ↪ training times
5 gc.collect()
6 with torch.no_grad():
7     torch.cuda.empty_cache()

```

Colab_Cell_4

```

1 %matplotlib inline
2 import matplotlib.gridspec as gridspec
3 import matplotlib.pyplot as plt
4 from google.colab import files
5
6 plt.rcParams['figure.dpi'] = 300
7 plt.rcParams['savefig.dpi'] = 300
8
9 # Function to plot the accuracy and loss graphs for the three configurations
  ↪ of LeNet
10 def plot():
11     trajectories = []
12     config = ['max', 'avg', 'no']
13     for POOLING in config:
14         with open(f'./results/results_{POOLING}.pkl', 'rb') as file:
15             trajectories.append(pickle.load(file))
16
17     idx = 0
18     def plot_trajectory(trajectory_name):
19         nonlocal idx, trajectories, config
20         plot_name = 'Loss' if trajectory_name == 'cross_entropy' else
           ↪ 'Accuracy'
21         gs = gridspec.GridSpec(4, 4)
22         m = 0

```

```

23
24     plt.figure(idx, figsize = (8,6))
25     for i in range(0, 4, 2):
26         for j in range(0, 4, 2):
27             if m == 3:
28                 break
29             if m < 2:
30                 ax = plt.subplot(gs[i:i+2, j:j+2])
31             else:
32                 ax = plt.subplot(gs[i:i+2, 1:3])
33
34             ax.plot(trajectories[m]['step'],
35                     trajectories[m][f'insample_{trajectory_name}'],
36                     linestyle='-.', color='b', label=f'Train {plot_name}')
37             ax.plot(trajectories[m]['step'],
38                     trajectories[m][f'outsample_{trajectory_name}'],
39                     linestyle='-', color='orange', label=f'Test {plot_name}')
40             ax.set_ylabel(plot_name)
41             ax.set_xlabel('Gradient Step')
42             ax.set_title(f'{config[m].upper()} pooling')
43             ax.legend(loc='best')
44
45             m+=1
46
47     plt.tight_layout(rect=[0, 0.03, 1, 0.95])
48     plt.suptitle(f'{plot_name} vs Gradient Steps')
49     plt.savefig(f'./plots/{plot_name}.png')
50     plt.show()
51     files.download(f'./plots/{plot_name}.png')
52     idx += 1
53
54     plot_trajectory('cross_entropy')
55     plot_trajectory('accuracy')
56
57 plot()

```

Colab_Cell_5

```

1 %matplotlib inline
2 import matplotlib.pyplot as plt
3 from google.colab import files
4
5 plt.rcParams['figure.dpi'] = 300

```

```

6 plt.rcParams['savefig.dpi'] = 300
7
8 # Function to plot the feature map visualizations for a given configuration
  ↪ of LeNet
9 def visualize(Pooling='avg'):
10     with open(f'./results/visualize_{Pooling}.pkl', 'rb') as file:
11         intermediates = pickle.load(file)
12
13     idx = 0
14     def visualize_images(out, out_name, nrows=2, ncols=3):
15         nonlocal idx
16         fig = plt.figure(idx, figsize=(ncols*2,nrows*2))
17         for fig_idx in range(1, nrows*ncols+1):
18             ax = fig.add_subplot(nrows, ncols, fig_idx)
19             ax.imshow(out[fig_idx-1], interpolation='nearest')
20             if out_name != 'input':
21                 ax.set_title(f'Channel {fig_idx}')
22             ax.axis('off')
23         plt.suptitle(f'Output Feature Maps of {out_name}' if out_name !=
  ↪ 'input' else 'Input Image')
24         plt.tight_layout(rect=[0, 0.03, 1, 0.95])
25         plt.savefig(f'./plots/visualize_{out_name}.png')
26         plt.show()
27         files.download(f'./plots/visualize_{out_name}.png')
28         idx += 1
29
30     print(intermediates['label'])
31     visualize_images(np.transpose(intermediates['input'].reshape((1, 3, 32,
  ↪ 32)), (0, 2, 3, 1)), 'input', 1, 1)
32     visualize_images(intermediates['conv1'].reshape((-1, 28, 28)), 'conv1',
  ↪ 2, 3)
33     visualize_images(intermediates['pool1'].reshape((-1, 14, 14)), 'pool1',
  ↪ 2, 3)
34
35 visualize('avg')

```